



Coalition for Content Provenance and Authenticity

C2PA Security Considerations

1.0, 2021-12-14: Release

Table of Contents

- 1. Introduction 2
 - 1.1. Threat modelling process overview 2
- 2. Security features 4
 - 2.1. Provenance model 4
 - 2.2. Trust model 5
 - 2.3. Claim signatures 5
 - 2.4. Content bindings 6
 - 2.5. Validation 6
 - 2.6. Protection of personal information 6
- 3. Security model 7
- 4. Threat model 9
 - 4.1. Threat and attack assumptions 9
 - 4.2. Threats 10



This work is licensed under a [Creative Commons Attribution 4.0 International License](https://creativecommons.org/licenses/by/4.0/).

Chapter 1. Introduction

This document provides information security considerations for technology described in the [C2PA core specification](#). It complements security-related content in other C2PA documents. This document includes:

- A summary of relevant security features of C2PA technology
- Security considerations for practical use of C2PA technology
- Threats to C2PA technology and respective treatment of those threats, including countermeasures

Analysis in this document does not assess security risk or likelihoods associated with threats or outcomes. These may vary based on a number of factors, such as the context in which C2PA technology is deployed and changes in the threat landscape. However, content in this document may be helpful in developing security risk and likelihood assessments related to use of C2PA technology.

IMPORTANT

Content in this document is non-normative and does not specify requirements for complying with the C2PA specification.

While the C2PA endeavours to develop and maintain a comprehensive threat model and security considerations for C2PA technology, we expect this work to be ongoing as technology develops and the threat landscape evolves.

IMPORTANT

This document is a work in progress.

1.1. Threat modelling process overview

The C2PA builds security into our designs as they are being developed, but also expects that security design and threat modelling will continue as the system, ecosystem, and threat landscape evolve.

To this end, the C2PA uses a focused threat modelling process to support development of a strong security and privacy design. Outcomes of the effort directly support development of documentation about explicit threats and security considerations while also facilitating security-oriented thinking throughout the design process.

The threat modelling process combines synchronous (i.e., live) threat modelling sessions with focused groups of Subject Matter Experts (SMEs) in the asynchronous development of content. The number of attendees in each synchronous session is kept small to promote efficient discussions, but all members of the C2PA have the opportunity to participate via either modality.

Like other security activities, we expect our threat modelling process to evolve with the C2PA ecosystem. Process documentation is considered a guide rather than a strict directive on how threat modelling works within the C2PA.

NOTE

In this document *asset* refers to C2PA assets and not [conventional threat model assets](#).

1.1.1. References

A variety of references and experiences are used to inform threat modelling and related security activities for the C2PA. This section provides a subset of public documents for reference.

- [IETF on security considerations](#)
- [IETF on privacy considerations \(guidelines\)](#)
- [W3C security and privacy self-review questionnaire](#)
- [OAuth2 threat model \(example\)](#)
- [Threat Modelling: Designing for Security](#)
- [OWASP Threat Modelling](#)
- [Microsoft Threat Modelling](#)

Chapter 2. Security features

This section summarizes C2PA features that support the [security model](#) specified later in this document. The [C2PA core specification](#) should be considered the canonical reference for C2PA features. It supersedes summaries provided here.

2.1. Provenance model

C2PA enables consumers to reason about asset provenance. C2PA specifies a model for asset provenance based on the following definitions.

Provenance

The logical concept of understanding the history of an *asset* and its interaction with *actors* and other *assets*, as represented by the *provenance data*.

Provenance data

The set of *manifests* for an *asset* and, in the case of a *composed asset*, its ingredients.

Assertion

A data structure which represents a statement asserted by an *actor* concerning the *asset*. This data is a part of the *manifest*.

Claim

A digitally signed and tamper-evident data structure that references a set of *assertions* by one or more *actors*, concerning an *asset*, and the information necessary to represent the *content binding*. If any *assertions* were redacted, then a declaration to that effect is included. This data is a part of the *manifest*.

Claim signature

The digital signature on the *claim* using the private key of an *actor*. The claim signature is a part of the *manifest*.

Manifest

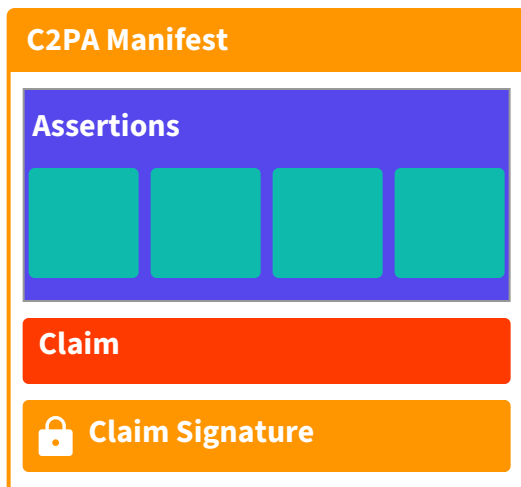
The set of information about the *provenance* of an *asset* based on the combination of one or more *assertions* (including *content bindings*), a single *claim*, and a *claim signature*. A manifest can either be embedded into an *asset* or be external to its *asset*.

Origin

The *manifest* in the *provenance data* which represents the method or device that initially created the *asset*.

Active Manifest

The last manifest in the list of manifests which is the one with the set of *content bindings* that are able to be validated.



2.2. Trust model

The basis of making trust decisions in C2PA, its Trust Model, is the identity of the actor associated with the cryptographic signing key used to sign the claim in the active manifest. The identity of a signatory is not necessarily a human actor, and the identity presented may be a pseudonym, completely anonymous, or pertain to a service or trusted hardware device with its own identity, including an application running inside such a service or trusted hardware. C2PA manifests can be validated indefinitely regardless of whether the cryptographic credentials used to sign its contents are later expired or revoked.

In addition to identity, C2PA leverages trusted timestamps to ensure the correctness of signing and revocation processes. A benefit of this approach is that revocation status can be optionally obtained when claims are generated (rather than when consumed), and revocation can be applied to subsets of signatures by timestamp. See the [Trust Model section of the specification](#) for more details.

2.3. Claim signatures

C2PA uses cryptographic signatures to ensure that any tampering of the asset provenance can be detected. C2PA does not offer any protection against the [complete removal of C2PA manifests from assets](#), but the design does aim to prevent any unauthorized tampering of either the digital asset content or the C2PA manifests themselves.

Details on supported cryptographic hashing and signature schemes, as well as format details, are covered in the [Cryptography section](#). As noted there, the C2PA maintains lists of supported algorithms and types that aim to adhere to cryptographic and compatibility best practices.

Detail on how signing credentials relate to identity are included in the [Trust Model section of the specification](#). The [Validation section](#) provides details on validating claims, which includes validating signatures.

2.4. Content bindings

Content bindings associate provenance data in C2PA manifests with assets.

In the context of a C2PA manifest, hard bindings provide a strict cryptographic binding between provenance and an asset; they aim to provide some useful guarantees against tampering. Soft bindings provide less strict binding between the provenance and an asset. They may associate C2PA provenance with multiple assets, and vice versa.

C2PA manifests may be embedded in assets or it may exist externally, such as when served as part of the same HTTP connection. In addition, C2PA manifests may also be decoupled from assets, enabling the C2PA manifest to be retrieved and processed separately from the asset.

Details on C2PA bindings are included in the [binding to content section of the C2PA core specification](#).

2.5. Validation

Validation checks provenance claim signatures and performs supporting operations. The validation algorithm is described in the [validation section of the C2PA specification](#).

2.6. Protection of personal information

C2PA provides features that can be used to protect confidentiality of personal information while still establishing the provenance of an asset.

- Anonymous or pseudo-anonymous identities may be used for signing claims.
- Assertions that may contain sensitive information can be removed via [redaction](#).
- Assertions can be customized for specific workflows, such as where specific information is required to remain encrypted end-to-end.
- Manifests may include [W3C Verifiable Credentials](#).

Note that confidentiality of personal information and other privacy-related concerns are addressed in the Harms Considerations document.

Chapter 3. Security model

The C2PA security model is focused on security of C2PA provenance data. While attackers may aim to compromise the security of other aspects of the system, this model assumes the pivotal goal is compromise of C2PA provenance.

The following principles should hold:

Integrity

Attempts to compromise the integrity of C2PA provenance data associated with a respective asset will be detected upon validation. This includes all data stored within C2PA manifests.

For embedded manifests, malicious association of an existing manifest with a differing asset is also tamper-evident. In addition, the following actions are tamper-evident for all manifests:

- Manifest: Malicious modification of an active or ingredient manifest.
- Claims and claim signatures: Addition, removal, or modification of an existing claim in an existing manifest.
- Assertions: Addition, removal, or modification of assertions in the assertion store referenced by an existing claim.

Availability

An attacker may degrade availability of C2PA data, but C2PA should provide means to mitigate this impact.

For example: An attacker may strip manifests from C2PA-enabled assets altogether. Guidance is provided on how to mitigate the impact of this scenario.

Confidentiality

C2PA does not specify confidentiality properties for assets themselves, provenance model primitives, private address books, or cryptographic secrets.

- They may be kept confidential via other means. For example by transmitting assets via TLS. However, enforcing this confidentiality is considered out-of-scope for this specification.
- Some private data, such as claim signing keys or private address books, must adhere to conventional confidentiality principles. Compromise of this private data could lead to indirect violation of other principles specified in this section. Enforcing confidentiality of this adjacent private data is considered out-of-scope for this specification.
- C2PA does not specify Digital Rights Management (DRM) technology.
- C2PA does not specify guarantees of confidentiality of personal information.
 - However, C2PA provides features that can be used to protect confidentiality of personal information ([summary](#))
 - Confidentiality of personal information and other privacy-related concerns are addressed in the TODO[Harms Considerations document].

In general, these assertions are achieved through application of the C2PA trust model to the claim generation process, and corresponding checks in the validation process.

Chapter 4. Threat model

This section documents information security threats to C2PA technology. Threats are described as threat scenarios and organized by categorically by prospective attacker goals. Where an attacker goal is what the attacker is trying to achieve, or the desired outcome of the threat scenario.

Each threat scenarios is described as follows:

- **Description:** A short description of the threat scenario
- **Prerequisites:** Relevant requirements for attempting an attack or threat scenario
- **Impact analysis:** Summary of the impact of the threat scenario on the C2PA security model
- **Guidance:** Any supplementary security guidance on addressing the threat scenario or treating potential residual security risk

Note that this document refers to threat sources, agents, actors and other malicious entities commonly as attackers.

4.1. Threat and attack assumptions

The following assumptions are made about entities that may attack C2PA technology:

- Attackers have full access to public assets and associated public metadata, including C2PA manifests.
- Attackers may attempt to access to claim generator software (e.g. C2PA-enabled photo editing software) as a typical user.
- Attackers do not have access to private keys referenced within the C2PA ecosystem (e.g. claim signing private keys, Timestamping Authority private keys, etc.). They may, however, attempt to access these keys via exploitation techniques.
- Attackers have the ability to publish assets to an impactful audience, e.g. as a social media user.
- Quantum computing cryptanalytic attacks are considered out-of-scope.

The following additional assumptions are made about the threat model:

- Attacks that rely on bypassing conventional security models of non-C2PA technology are considered out-of-scope. For example:
 - Attacks against the C2PA Validator as an architectural component within a user agent (e.g. web browser) are considered in-scope.
 - A compromise of the C2PA validation logic via some non-C2PA attack vector within the user agent is considered out-of-scope, since this would typically require bypassing the security model of the user agent (e.g. escaping a web browser sandbox or compromising the host via other means).
- Attacks across conventional trust boundaries are considered in-scope.

- For example, attacks against C2PA-enabled assets traversing the internet.

4.2. Threats

4.2.1. Goal: Degrading access to provenance information

Threat scenarios in this category aim to prevent media consumers from reasoning about the provenance of assets. In terms of the security model, this is a compromise of availability of C2PA provenance data.

4.2.1.1. Threat: Stripping C2PA manifests

4.2.1.1.1. Description

An attacker removes C2PA manifests from an asset delivered with hard bindings. For example, by downloading C2PA-enabled media from a social media site, stripping the C2PA manifest(s), and re-posting it.

4.2.1.1.2. Prerequisites

Technical ability to strip provenance data from published or stored assets and deliver the stripped assets targeted users or systems. For example, via republishing to social media or on-path attacks.

4.2.1.1.3. Impact analysis

It is possible for an attacker to remove metadata as described. Depending on the context, this could cause a user agent not to display provenance information associated with the asset.

4.2.1.1.4. Security guidance

The impact of this threat scenario could be mitigated through use of decoupled manifests, e.g. by making decoupled C2PA manifests available to user agents when embedded manifests could be missing.

Use of decoupled manifests can introduce security and privacy issues. See the [Content binding security section](#) for more information.

4.2.1.2. Threat: Stripping from data within C2PA manifests

4.2.1.2.1. Description

An attacker removes metadata from within C2PA manifests(s) stored in a asset delivered with hard bindings, such as the assertion store, individual assertions, claims, claim signatures, etc. For example, by downloading C2PA-enabled media from a social media site, stripping the target C2PA metadata, and re-posting it.

4.2.1.2.2. Prerequisites

Technical ability to strip provenance data from published or stored assets and deliver the stripped assets targeted users or systems. For example, via republishing to social media or on-path attacks.

4.2.1.2.3. Impact analysis

It is possible for an attacker to remove metadata as described. However, [validation checks](#) (e.g. as implemented by a user agent) will detect removal or partial removal of metadata stored within C2PA manifests. Removal of an entire C2PA manifest is possible, and described as a [distinct threat scenario](#).

This is similar to the case of stripping all C2PA metadata. However, in this scenario there is an added benefit that user agents can explicitly detect the attempt to remove metadata as part of the validation process.

4.2.1.2.4. Security guidance

The impact of this threat scenario could be mitigated through explicit detection of missing metadata as described here, and factoring this into system behaviour. For example, a user agent may modify display of the affected asset.

The impact of this threat scenario could be also mitigated through use of decoupled manifests, e.g. by making decoupled C2PA manifests available to user agents when embedded metadata is missing.

Use of decoupled manifests can introduce security and privacy issues. See the [Content binding security section](#) for more information.

4.2.2. Goal: Tamper with provenance information

Threat scenarios in this category aim to modify or spoof provenance information associated with C2PA-enabled assets, e.g. to present misleading provenance data to media consumers. This is considered a compromise of integrity of provenance data in terms of the C2PA security model.

4.2.2.1. Threat: Copying C2PA metadata

4.2.2.1.1. Description

An attacker copies valid, hard-bound C2PA metadata, such as assertions, claims, or entire C2PA manifests, from one C2PA-enabled asset to another.

4.2.2.1.2. Prerequisites

Technical ability to strip provenance data from published or stored assets and deliver the stripped assets targeted users or systems. For example, via republishing to social media or on-path attacks.

4.2.2.1.3. Impact analysis

An attacker is able to perform the copying as described. However, C2PA hard bindings provide a cryptographic binding between assets and C2PA manifests via claim signatures. Thus, [validation checks](#) (e.g. as implemented by a user agent) will detect the replaced C2PA manifests.

This is similar to the case of stripping all C2PA metadata. However, in this scenario there is an added benefit that user agents can explicitly detect the attempt to remove metadata as part of the validation process.

4.2.2.1.4. Security guidance

The impact of this threat scenario could be mitigated through explicit detection of incorrect metadata as described here, and factoring this into system behaviour. For example, a user agent may modify display of the affected asset.

The impact of this threat scenario could also be mitigated through use of decoupled manifests, e.g. by making decoupled C2PA manifests available to user agents when metadata fails validation.

Use of decoupled manifests can introduce security and privacy issues. See the [Content binding security section](#) for more information.

4.2.2.2. Threat: Spoofing signed C2PA metadata via stolen key

4.2.2.2.1. Description

An attacker compromises the confidentiality of a claim signing key and uses it to attach valid, malicious provenance to a asset.

4.2.2.2.2. Prerequisites

Ability to read and use a valid C2PA claim signing key. For example by stealing it from a claim generation and signing system. Ability to copy, modify, or generate assets, add desired (valid) C2PA provenance metadata to them, and deliver them to targeted users or systems. For example, via republishing to social media or on-path attacks.

4.2.2.2.3. Impact analysis

An attacker that is able to steal a signing key can spoof valid C2PA manifests on arbitrary assets. This applies to both hard and soft bindings. Valid claim signatures will be limited to the exposed key; an attacker cannot impersonate other signers in this scenario. These manifests will pass validation checks. Further, an attacker will be able to continue to generate assets with spoofed provenance until the compromised credential is revoked, and revocation status is propagated to the targeted users/systems.

If the signer's credential type does not support or provide a means for revocation, it may be necessary to remove the trust anchor associated with the signer's credential from the respective C2PA trust list to effectively revoke the malicious manifests. In this case, all manifests associated with the trust anchor will be invalidated.

Revocation of the signing key associated with the malicious manifests may invalidate other manifests if signing keys are reused between claim generation / claim signing system users.

4.2.2.2.4. Security guidance

Claim generators must adhere to best practices for securing keys in their respective operating environments. This includes using platform-specific best practices for key storage and management (e.g. hardware security modules, key management services), minimizing key reuse, and revoking keys when compromise is suspected. For more information on key management, see the [NIST Key Management Guidelines](#).

It is highly recommended that all signing credentials include revocation information.

4.2.2.3. Threat: Spoofing signed C2PA metadata via misuse of claim generator

4.2.2.3.1. Description

An attacker misuses a legitimate claim generator (e.g. photo editing service) to add misleading provenance to a C2PA-enabled asset.

4.2.2.3.2. Prerequisites

Ability to access and exploit or misuse a claim generator system. For example, misuse of a claim generator may result in malicious additions of C2PA manifests or Claims to assets; exploitation of vulnerabilities in claim generators may lead to modification or tampering with existing C2PA manifest data added with the target claim generator's signing key(s). Ability to deliver assets to targeted users. For example, via publishing to social media or on-path attacks.

4.2.2.3.3. Impact analysis

An attacker may be able to misuse or exploit a claim signing system to create malicious C2PA manifests. These can then be delivered to targeted users/systems. Note that valid claim signatures will be limited to credentials used by the targeted claim signing system; an attacker cannot impersonate other signers in this scenario.

If the signer's credential type does not support or provide a means for revocation, it may be necessary to remove the trust anchor associated with the signer's credential from the respective C2PA trust list to effectively revoke the malicious manifests. In this case, all manifests associated with the trust anchor will be invalidated.

Revocation of the signing key associated with the malicious manifests may invalidate other manifests if signing keys are reused between claim generation / claim signing system users.

4.2.2.3.4. Security guidance

Claim generators must adhere to industry best practices for information security and anti-abuse practices, including leveraging available platform-specific features for deployment (e.g. [Android SafetyNet](#), [Apple DeviceCheck](#)). Likewise, claim generators should use platform-specific best practices for key storage and management. For example, usage of hardware security modules, key management services), minimizing key reuse, and revoking keys when compromise is suspected. For more information on key management, see the [NIST Key Management Guidelines](#).

It is highly recommended that all signing credentials include revocation information.

The impact of this threat scenario could be mitigated through use of decoupled manifests, e.g. by making decoupled C2PA manifests available to user agents when a user suspects falsified provenance data.

Use of decoupled manifests can introduce security and privacy issues. See the [Content binding security section](#) for more information.

4.2.2.4. Threat: Issuance of claim signing keys to malicious entity

4.2.2.4.1. Description

An attacker causes a C2PA claim signing key issuer to issue a valid signing key to them. They then use this signing key to add misleading provenance to a C2PA-enabled media asset.

4.2.2.4.2. Prerequisites

Ability to bypass unspecified processes for obtaining a claim signing key.

4.2.2.4.3. Impact analysis

An attacker that is able to obtain a legitimate claim signing key can spoof valid C2PA manifests on arbitrary media assets. This applies to both hard and soft bindings. These manifests will pass validation checks. Further, an attacker will be able to continue to generate media assets with spoofed provenance until the compromised credential is revoked, and revocation status is propagated to the targeted users/systems.

It may be necessary to remove the trust anchor associated with the signer's credential from the respective C2PA trust list to effectively revoke malicious manifests. In this case, all manifests associated with the trust anchor will be invalidated.

4.2.2.4.4. Security guidance

The C2PA does not currently specify requirements regarding issuance of claim signing keys. It is recommended that these processes be robust against exploitation and misuse. The [CAB Forum Baseline Requirements](#) provide a working example of security requirements for signing credential issuers.

4.2.2.5. Threat: Swapping ingredients

4.2.2.5.1. Description

An attacker finds a media asset with a C2PA manifest block that includes an ingredient manifest with a signature from valid, non-public credentials. They replace the ingredient manifest with a malicious manifest.

4.2.2.5.2. Prerequisites

Ability to discover media with susceptible ingredient manifest and generate a malicious replacement. And deliver the media asset to the target users/systems.

4.2.2.5.3. Impact analysis

Validation of ingredients is optional. See the [Validation section of the C2PA specification](#) for details.

A ingredient validator that does not have access to the original, non-public signing credential cannot differentiate between the original and compromised provenance. Both ingredient manifests will fail validation. An ingredient validator that does have access to the original signing credential will be able to differentiate between the two C2PA-enabled media assets. The original manifest will validate, and the malicious manifest will fail validation.

4.2.2.5.4. Security guidance

Validators that apply private address books, and also act as claim generators, can optionally use C2PA reviews (see [Use of Schema.org](#)) to propagate their assessment of ingredients to downstream consumers. When this validation is performed on both media assets described in this scenario, downstream consumers will be able to differentiate between them.

4.2.2.6. Threat: Binding malicious media asset to valid provenance via soft-binding collision

4.2.2.6.1. Description

An attacker generates a malicious media asset whose soft binding matches the soft binding of a legitimate, target media asset. The attacker creates a decoupled, soft-bound manifest for the malicious media asset. The malicious media asset is then delivered to the target/system user, which retrieves the legitimate C2PA manifest via the common soft binding.

4.2.2.6.2. Prerequisites

Ability to generate a colliding soft binding for a target media asset and publish the resulting manifest such that it will be retrieved by the target user/system for the target media asset.

4.2.2.6.3. Impact analysis

Malicious media bound to legitimate C2PA provenance data is validated and presented to the target system/user.

4.2.2.6.4. Security guidance

C2PA provides guidance on applying soft bindings, which may include algorithms for mitigating this type of attack. See the [content binding documentation](#) for details.

4.2.2.7. Threat: Use valid W3c Verified Credentials to add authenticity to incorrect actor attribution

4.2.2.7.1. Description

An attacker copies a valid W3C Verified Credential (VC) from an existing C2PA manifest. The attacker then creates a separate, valid C2PA manifest that includes the copied VC, misleading information, and a valid claim signature. The resulting manifest is presented to consumers.

4.2.2.7.2. Prerequisites

- Ability to copy a valid VC from an existing C2PA manifest
- Ability to generate and sign a claim that passes validation
- Ability to deliver assets to consumer(s), for example via publishing to social media

4.2.2.7.3. Impact analysis

Consumers are presented with UX that indicates the subject of copied VC is associated with the manipulated asset.

4.2.2.7.4. Security guidance

User agents must make a clear separation between the signer of the claim's assertions and any included W3C Verifiable Credentials. It must be clear that the association of those VCs to assertions, are made by the claim signer (i.e. the subject of the VC was not responsible for the association).

Claim generators and signing operations should adhere to security best practices.

4.2.2.8. Threat: Binding valid media asset to malicious provenance via soft-binding collision

4.2.2.8.1. Description

An attacker generates a dummy media asset whose soft binding matches the soft binding of another media asset. They then use a claim generator to create a valid, malicious C2PA claim(s), and a decoupled C2PA manifest. They then publish the C2PA manifest. When the target media asset is validated, the malicious C2PA manifest is retrieved via the common soft binding.

4.2.2.8.2. Prerequisites

Ability to generate a colliding soft binding for a target media asset and publish the resulting manifest such that it will be retrieved by the target user/system for the target media asset.

4.2.2.8.3. Impact analysis

Malicious provenance data bound to a legitimate media asset is validated and presented to the target system/user.

4.2.2.8.4. Security guidance

C2PA provides guidance on applying soft bindings, which may include algorithms for mitigating this type of attack. See the [content binding documentation](#) for details.

4.2.3. Goal: Exploit functionality introduced by C2PA

Threat scenarios in this category do not aim to break the C2PA security model. Rather, they aim to leverage C2PA functionality to perform systems attacks or exploitation. Note that the set of scenarios included here is not meant to be comprehensive; they focus on application of known exploitation techniques against behaviour explicitly referenced in the C2PA specification.

(This section is to be expanded)